

IST387_FinalProj

Nicholas Santangelo

2022-12-05

```
#1
library(tidyverse) #library calls the tidyverse package

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr 0.3.4
## v tibble 3.1.8       v dplyr 1.0.10
## v tidyr 1.2.1        v stringr 1.4.1
## v readr 2.1.2        v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

costDF <- read_csv('https://intro-datascience.s3.us-east-2.amazonaws.com/HMO_data.csv')

## Rows: 7582 Columns: 14
## -- Column specification -----
## Delimiter: ","
## chr (8): smoker, location, location_type, education_level, yearly_physical, ...
## dbl (6): X, age, bmi, children, hypertension, cost
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

##read_csv function takes data from the linked file. Placed into costDF variable.

str(costDF) #str function gets the overall structure of the costDF df

## spec_tbl_df [7,582 x 14] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ X : num [1:7582] 1 2 3 4 5 7 9 10 11 12 ...
## $ age : num [1:7582] 18 19 27 34 32 47 36 59 24 61 ...
## $ bmi : num [1:7582] 27.9 33.8 33 22.7 28.9 ...
## $ children : num [1:7582] 0 1 3 0 0 1 2 0 0 0 ...
## $ smoker : chr [1:7582] "yes" "no" "no" "no" ...
## $ location : chr [1:7582] "CONNECTICUT" "RHODE ISLAND" "MASSACHUSETTS" "PENNSYLVANIA" ...
## $ location_type : chr [1:7582] "Urban" "Urban" "Urban" "Country" ...
## $ education_level: chr [1:7582] "Bachelor" "Bachelor" "Master" "Master" ...
## $ yearly_physical: chr [1:7582] "No" "No" "No" "No" ...
## $ exercise : chr [1:7582] "Active" "Not-Active" "Active" "Not-Active" ...
## $ married : chr [1:7582] "Married" "Married" "Married" "Married" ...
## $ hypertension : num [1:7582] 0 0 0 1 0 0 0 1 0 0 ...
```

```
## $ gender      : chr [1:7582] "female" "male" "male" "male" ...
## $ cost        : num [1:7582] 1746 602 576 5562 836 ...
## - attr(*, "spec")=
## .. cols(
## ..   X = col_double(),
## ..   age = col_double(),
## ..   bmi = col_double(),
## ..   children = col_double(),
## ..   smoker = col_character(),
## ..   location = col_character(),
## ..   location_type = col_character(),
## ..   education_level = col_character(),
## ..   yearly_physical = col_character(),
## ..   exercise = col_character(),
## ..   married = col_character(),
## ..   hypertension = col_double(),
## ..   gender = col_character(),
## ..   cost = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

*#The df has 14 variables with 7,582 observations. All columns are either a number
#or character data type. DF appears to represent personal health care costs*

summary(costDF) #summary function produces descriptive stats for costDF columns

```
##           X                age                bmi                children
## Min.      :      1  Min.    :18.00  Min.    :15.96  Min.    :0.000
## 1st Qu.:    5635  1st Qu.:26.00  1st Qu.:26.60  1st Qu.:0.000
## Median :   24916  Median :39.00  Median :30.50  Median :1.000
## Mean   :  712602  Mean   :38.89  Mean   :30.80  Mean   :1.109
## 3rd Qu.:  118486  3rd Qu.:51.00  3rd Qu.:34.77  3rd Qu.:2.000
## Max.    :13110111  Max.    :66.00  Max.    :53.13  Max.    :5.000
##                                     NA's      :78
##           smoker           location           location_type           education_level
## Length:7582           Length:7582           Length:7582           Length:7582
## Class :character      Class :character      Class :character      Class :character
## Mode  :character      Mode  :character      Mode  :character      Mode  :character
##
##
##
## yearly_physical           exercise           married           hypertension
## Length:7582           Length:7582           Length:7582           Min.    :0.0000
## Class :character      Class :character      Class :character      1st Qu.:0.0000
## Mode  :character      Mode  :character      Mode  :character      Median :0.0000
##                                     Mean   :0.2005
##                                     3rd Qu.:0.0000
##                                     Max.    :1.0000
##                                     NA's    :80
##           gender           cost
## Length:7582           Min.    :      2
## Class :character      1st Qu.:   970
## Mode  :character      Median : 2500
```

```
##           Mean    : 4043
##           3rd Qu.: 4775
##           Max.    :55715
##
```

#cost df appears to have: an age range of 18-66, a child range of 0-5, and an overall cost range of \$2-\$55,715

```
#2
names(which(colSums(is.na(costDF))>0)) #is.na checks for null values in the
```

```
## [1] "bmi"           "hypertension"
```

#specified costDF Dataframe. colSums gets the sum count of the null values, if greater than 0. Which gets the column number of columns with null values. Names of function produces the columns name's

```
library(imputeTS) #library calls the imputeTS package
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```

```
typeof(costDF$bmi) #typeof function gets the data type of the bmi column
```

```
## [1] "double"
```

```
typeof(costDF$hypertension) #typeof function gets the data type of the
```

```
## [1] "double"
```

#hypertension column

```
costDF$bmi <- na_interpolation(costDF$bmi) #redefines bmi column using the
#na_interpolation function from the imputeTS package. Gets calculated replacement
#values for null rows. bmi column defined as data to be used.
```

```
costDF$hypertension <- na_interpolation(costDF$hypertension) #redefines bmi
#column using the na_interpolation function from the imputeTS package. Gets
#calculated replacement values for null rows. bmi column defined as data to be
#used.
```

```
sum(is.na(costDF$bmi)) #sum function used to callculated the total of null
```

```
## [1] 0
```

#values found using the is.na function. bmi column defined as data to be checked

```
sum(is.na(costDF$hypertension)) #sum function used to callculated the total of
```

```
## [1] 0
```

```
#null values found using the is.na function. hypertension column defined as data  
#to be checked
```

```
#3  
exercercise_table <- table(costDF$exercercise) #table function used to generate  
#breakdown of column values. exercercise defined as column to be used. put into  
#variable exercercise_table  
exercercise_table #displays the table
```

```
##  
##      Active Not-Active  
##      1888      5694
```

```
#around one-third of people are active and two-thirds are not.
```

```
locationT_table <- table(costDF$location_type) #table function used to generate  
#breakdown of column values. location_type defined as column to be used.put into  
#variable locationT_table  
locationT_table #displays the table
```

```
##  
## Country      Urban  
##      1903      5679
```

```
#around 25% of the people live in the country and 75% live in an urban area
```

```
education_table <- table(costDF$education_level) #table function used to generate  
#breakdown of column values. education_level defined as column to be used.put  
#into variable education_table  
education_table #displays the table
```

```
##  
##      Bachelor      Master No College Degree      PhD  
##      4578      1533      759      712
```

```
#the majority hold a bachelor, then masters is second most common, followed by  
#no degree, then PhD.
```

```
#4  
summary(costDF$cost)[5] #summary function displays the quantitative stats on the
```

```
## 3rd Qu.  
##      4775
```

```
#defined cost column data. Subset used to get just the third quartile number
```

```
costDF$expensive <- ifelse(costDF$cost > summary(costDF$cost)[5], 'Expensive',  
                           'Inexpensive')  
#expensive column defined. ifelse used to create a boolean expression. if cost  
#is bigger than the 75% value than it would be 'Expensive', else it would be
```

```
#'Inexpensive'
```


```
#I decided that the 3rd quartile or the 75% mark would be the threshold because  
#the mean is significantly bigger than the median, meaning there are a large  
#number of values bigger than the median. With that said, I felt like the top  
#25% of values would be a good mark for the expensive threshold.
```

```
#5
```

```
exp_df <- subset(costDF, expensive == 'Expensive') #exp_df defined as a dataframe.  
#subset function used to take a portion of the defined costDF data. Expensive  
#column used to subset the data, only selects rows where 'Expensive' is the  
#expensive column's value  
inexp_df <- subset(costDF, expensive == 'Inexpensive') #inexp_df defined as a  
#dataframe.subset function used to take a portion of the defined costDF data.  
#Expensive column used to subset the data, only selects rows where 'Inexpensive'  
#is the expensive column's value
```

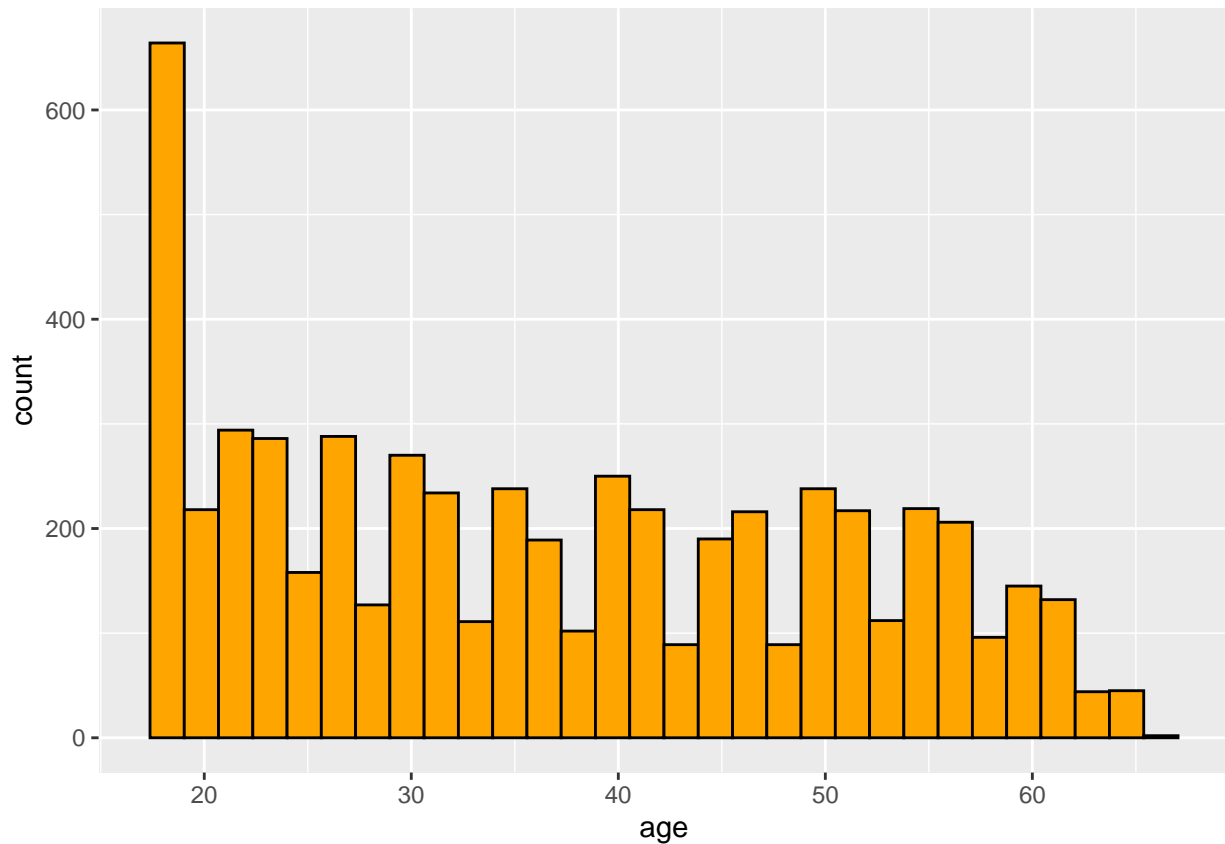
```
ggplot(exp_df, aes(x=age)) +geom_histogram(fill = "light Blue", color = 'black')
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



```
#ggplot used to map the exp_df defined data. aesthetics defined as x axis value  
#equals age. geom_histogram used to identify the graph type. "light Blue" defined  
#as the fill color and 'black' defined as the outline color.  
ggplot(inexp_df, aes(x=age)) +geom_histogram(fill = "orange", color = 'black')
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




```
#ggplot used to map the inexp_df defined data. aesthetics defined as x axis value  
#equals age. geom_histogram used to identify the graph type. "orange" defined  
#as the fill color and 'black' defined as the outline color
```

```
#expensive graph appears to be right leaning, indicating that those that are  
#expensive are likely to be older. Inexpensive graph is left leaning indicating  
#that younger people are likely to be cheaper.
```

```
#9 continued
```


```
ggplot(exp_df, aes(x=bmi)) +geom_histogram(fill = "light Blue", color = 'black')
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



```
#ggplot used to map the exp_df defined data. aesthetics defined as x axis value  
#equals bmi geom_histogram used to identify the graph type. "light Blue" defined  
#as the fill color and 'black' defined as the outline color.  
ggplot(inexp_df, aes(x=bmi)) +geom_histogram(fill = "orange", color = 'black')
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




*#ggplot used to map the inexp_df defined data. aesthetics defined as x axis value
#equals bmi geom_histogram used to identify the graph type. "orange" defined
#as the fill color and 'black' defined as the outline color.*

*#while both graphs do show a relatively similar standard deviation, the peak
#count of the expensive graph is on a higher bmi (~35) than that of the
#inexpensive graph (~27)*

#9 continued


```
ggplot(exp_df, aes(x=cost))+geom_histogram(fill = "light Blue", color = 'black')
```

'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.



```
#ggplot used to map the exp_df defined data. aesthetics defined as x axis value  
#equals cost geom_histogram used to identify the graph type. "light Blue" defined  
#as the fill color and 'black' defined as the outline color.  
ggplot(inexp_df, aes(x=cost))+geom_histogram(fill = "orange", color = 'black')
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



*#ggplot used to map the inexp_df defined data. aesthetics defined as x axis value
#equals cost geom_histogram used to identify the graph type. "orange" defined
#as the fill color and 'black' defined as the outline color.*

*#the expensive histogram shows a heavily left leaning, steep distribution of cost
#where as the inexpensive is also left leaning, just much more evenly distributed.
#This distribution implies that there are a lot of people of the 3rd quartile
#threshold that I decided, meaning I may have set the expensive boundary a
#little too early, maybe 80-90% would have been better.*

```
#6
state <- data.frame(
  cost = aggregate(costDF$cost, list(costDF$location), mean))
#state defined as a variable name, data.frame used to create this as a dataframe.
#cost identified as column name, uses aggregate function to get mean cost of
#cost column based on the different location values.

colnames(state)[1] <- "name"
#renames location values column as "name"
colnames(state)[2] <- "ave_cost"
#renames average cost values column as "ave_cost"
```

```
#7
state[which.max(state$ave_cost),]
```

```
##      name ave_cost
## 5 NEW YORK 4661.506
```

```
#state df subseted to output a singular row. which.max function used to
#determine which row has the highest value. ave_cost defined as the data used.
```


```
#8
#install.packages('usmap')
#install.packages packages function used to install the usmap package
library(usmap)
#library function used to call and load in the usmap package
library(ggplot2)
#library function used to call and load in the ggplot2 package

state$state <- c("CT", "MD", "MA", "NJ", "NY", "PA", "RI")
#state column defined for state df. Abbreviations of state names inputed as data
#so states can be mapped

plot_usmap(data= state,
            values = "ave_cost",
            include= c("CT", "MD", "NJ", "NY", "PA", "RI", "MA")) +
  scale_fill_continuous(low= 'white', high = 'orange', name= 'Average Cost', label = scales::c
  ggtitle("Average Cost per State") + theme(plot.title = element_text(size=30))
```

```
## Warning: Ignoring unknown parameters: linewidth
```

Average Cost per State



*#plot_usmap function used to plot the map of the US. state df defined as the
#data to be used for color gradient. ave_cost defined as column to base gradient
#on. unclude parameter used to indicate which states I want to appear in the map.
#scale_fill_continuous added to define the mapping colors as well as the title
#to the legend. ggtitle added to create a overall graph title.*

*#based on the color gradient, New York is the most expensive. Then Massachusetts
#is the second most expensive. After that, Rhode Island and Pennsylvania are
#about equal for third place. Then, New Jersey and Connecticut are about equal
#for fourth place. Lastly is Maryland which is the cheapest.*

#9

```
costCat <- data.frame(location=as.factor(costDF$location),  
                      location_type=as.factor(costDF$location_type),  
                      expensive =as.factor(costDF$expensive),  
                      education_level =as.factor(costDF$education_level),  
                      yearly_physical =as.factor(costDF$yearly_physical),  
                      gender = as.factor(costDF$gender),  
                      hypertension = as.factor(costDF$hypertension),  
                      exercise = as.factor(costDF$exercise),  
                      age= costDF$age,  
                      bmi= costDF$bmi,  
                      x= costDF$x,  
                      married = as.factor(costDF$married),  
                      smoker= as.factor(costDF$smoker))
```

*#costCat variable defined. data.frame function used to create a dataframe. Almosts
#all variables from costDF function inputted. Those that have binary-type or
#select number of possible responses, set as a factor type. Those that have a
#wide range of values (age, bmi, and x) simply converted over in their original
#form. Cost column was excluded as we are trying to find what variables
#determine the cost.*

```
library(arules) #library calls and loads in the arules package
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
##      expand, pack, unpack
```

```
##
```

```
## Attaching package: 'arules'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      recode
```

```
## The following objects are masked from 'package:base':
##
##   abbreviate, write
```


```
library(arulesViz) #library calls and loads in the arulesViz package
```

```
costCat_1 <- as(costCat, 'transactions')
```

```
## Warning: Column(s) 9, 10, 11 not logical or factor. Applying default
## discretization (see '? discretizeDF').
```

```
#costCat_1 defined as the variable name. as function used to convert costCat
#dataframe to a transaction class data type. We do this to make the data able
#for further mining and rule identification.
```

```
itemFrequencyPlot(costCat_1, topN=15)
```



```
#itemFrequencyPlot function used to get the most common single column value
#occurrences. costCat_1 defined as the data to be used.topN paramater function
#used to only select the tope 15 most common values.
```

```
#looks like 80% of the people are not smokers, and roughly 80% of people don't
#have hyper tension. And arguably because of this, around 75% of people are
#inexpensive.
```

```
#10
cost_analysis <- apriori(costCat_1,
                        parameter=list(supp=0.1, conf=0.7),
                        control=list(verbose=F),
                        appearance=list(default="lhs",rhs= 'expensive=Expensive'))

#cost_analysis defined as the variable name. apriori function used to generate
#association rules based on the costCat_1 value occurrence data from the previous
#question. parameter attribute used to designate required support level of 10%
#and confidence level of 70%, this was to narrow down strongest correlations.
#defines control as a list with the verbose flow type as F. appearance is
#defined as a list with the default labeled as lhs. the rhs variable is
#identified as expensive=Expensive

inspect(cost_analysis)
```

	lhs	rhs	support	confidence	coverage	lift	count
## [1]	{smoker=yes}	=> {expensive=Expensive}	0.1424426	0.7302231	0.1950673	2.921663	1080
## [2]	{location_type=Urban, smoker=yes}	=> {expensive=Expensive}	0.1067001	0.7262118	0.1469269	2.905614	809
## [3]	{exercise=Not-Active, smoker=yes}	=> {expensive=Expensive}	0.1168557	0.8158379	0.1432340	3.264213	886
## [4]	{yearly_physical=No, smoker=yes}	=> {expensive=Expensive}	0.1065682	0.7169476	0.1486415	2.868547	808
## [5]	{hypertension=0, smoker=yes}	=> {expensive=Expensive}	0.1109206	0.7237522	0.1532577	2.895772	841

*#inspect function displays the results of cost_analysis showing the
#frequency and association of various attribute combinations.*

*#not surprising, the most common value association with expensive people is
#smoking. With a confidence of 73%, it means that nearly 3 out of 4 expensive
#people smoke. The second greatest association was smokers living in an urban
#environment with a confidence slightly lower than just smokers alone. In third
#was smokers who did not exercise, which had the highest confidence of 81.5%.
#This means you are more likely to be expensive if you do both than if you just
#smoked. Because this pairing isn't as common (support) it only ranks third on
#highest association.*

*#support: The percentage this pairing occurs in the whole dataset (out of
#expensive people)*

*#confidence: The percentage that if you have the lhs characteristics then you
#will have the rhs.*

*#lift: Shows the effectiveness of the model, meaning it is the ratio between the
#confidence of the rule and the expected confidence. Higher lift indicates a
#higher correlation.*

```
#11
library(caret) #library calls and loads in the caret package
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
cost_train <- createDataPartition(y=costDF$expensive,p=.70,list=FALSE)
#cost_train assigned as variable name. createDataPartition function used to get
#a sample of the data. expensive column defined as the data to be sampled. P set
#at .7 to sample 70% of the data. list attribute set as false.
train_set <- costDF[cost_train,]
#train_set defined of subset of costDF dataframe to only include rows from the
#cost_train data partition
test_set <- costDF[-cost_train,]
#test_set defined of subset of costDF dataframe to exclude all rows from the
#cost_train data partition

dim(test_set)[1] + dim(train_set)[1]
```

```
## [1] 7582
```

```
#calculates the total observations between the test_set and train_set. dim
#used to get the dimensions, subset used to isolate the observations
dim(costDF)[1]
```

```
## [1] 7582
```

```
#dim function used on the costDF dataframe to check if the number of
#observations match. subset used to isolate the observations
```

```
#12
```

```
library(kernlab) #library calls and loads in the kernlab package
```

```
##
```

```
## Attaching package: 'kernlab'
```

```
## The following object is masked from 'package:arules':
```

```
##
```

```
## size
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## cross
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## alpha
```

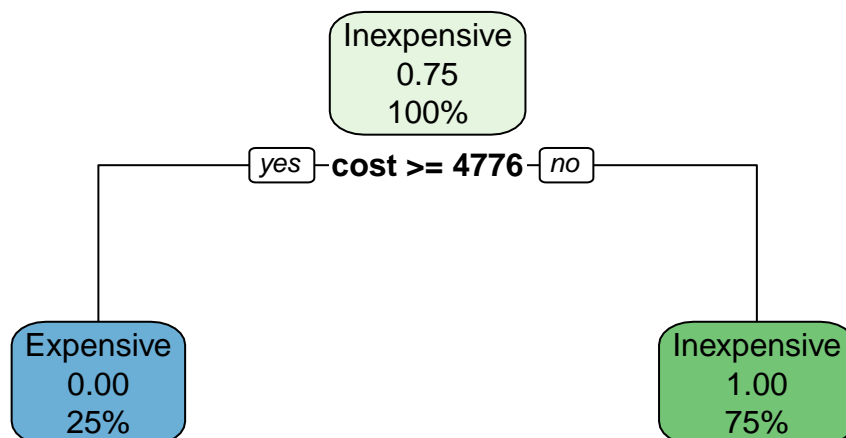
```

library(rpart) #library calls and loads in the rpart package
library(rpart.plot) #library calls and loads in the rpart.plot package
library(caret) #library calls and loads in the caret package

model <- ksvm(as.factor(expensive)~ .,data=train_set, C =5, cross =3,
              prob.model = TRUE)
#model defined as variable name. ksvm function called from kernlab package to
#create a support vector. Defines expensive as dependant variable, defines
#train_set as data, defines count as 5, defines cross as 3, defines prob.model
#as true to builds a model for calculating probabilities

cost_tree <- rpart(expensive~ .,data=train_set)
#cost_tree assigned as variable name. uses the rpart function to develop a data
#map of the expensive column in the train_set df
rpart.plot(cost_tree)

```



#rpart.plot function used to plot the cost_tree rpart of the expensive variable

#12 Continued

```

cost_predict <- predict(model, test_set, type = 'response') #uses predict
#function to test test data, defines type as response

confusion_table <- table(cost_predict, test_set$expensive)
#Creates table of prediction's data, identifies cost_predict and expensive as
#the data. The 4 numbers of the table represent the count of expensive and

```



```
#inexpensive of the model. The diagnosis of the table are the corespondent and
#prediction of the model.
confusion_table
```

```
##
## cost_predict  Expensive Inexpensive
##   Expensive      550        14
##   Inexpensive     18       1692
```

```
#displays the table confusion_table
```

```
error <- (sum(confusion_table) - sum(diag(confusion_table))) / sum(confusion_table)
#Gets the sum of table subtracted by the sum of the diagonal, devided by the
#total sum to get the error rate
error
```

```
## [1] 0.01407212
```

```
#Displays the error rate
```

```
#12 Continued
confusion_matrix <- confusionMatrix(cost_predict, as.factor(test_set$expensive))
#confusionMatrix function gets the calculation of observed and predicted classes.
#cost_predict defined as the prediction used, expensive defined as the column to
#be examined.
confusion_matrix
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction  Expensive Inexpensive
##   Expensive      550        14
##   Inexpensive     18       1692
##
##               Accuracy : 0.9859
##               95% CI : (0.9802, 0.9904)
##   No Information Rate : 0.7502
##   P-Value [Acc > NIR] : <2e-16
##
##               Kappa : 0.9624
##
## Mcnemar's Test P-Value : 0.5959
##
##               Sensitivity : 0.9683
##               Specificity : 0.9918
##   Pos Pred Value : 0.9752
##   Neg Pred Value : 0.9895
##   Prevalence : 0.2498
##   Detection Rate : 0.2419
##   Detection Prevalence : 0.2480
##   Balanced Accuracy : 0.9801
##
```

```
##          'Positive' Class : Expensive
##
```

```
#displays the confusion matrix
```

```
#13
```

```
#Based on the nearly identical error rates of 1.18% and 1.19%, It would be  
#incorrect of me to chose one model over the other. With their incredible  
#accuracy, both of these models are more than adequate.
```

```
#14
```

```
#Hello CEO,
```

```
#after a lot of examination and analysis, there are a few takeaways regarding  
#what variables effect a persons insurance cost the most. First and most  
#importantly, smoking is something that has a massive impact on a person's cost.  
#With a confidence percentage of 73%, smoking is the single greatest contributor  
#to expensive cost. Beyond smoking, factors like lack of exercise and living in  
#an urban environment can also play a pretty important role. Beyond factors like  
#these, we've discovered that the a person's age can be an indication of their  
#cost, with the majority of our expensive clients over the age of 45. As far as  
#the expectations for next year go, I would expect your out-of-shape elderly  
#smokers to be your most expensive clients. But, out-of-shape smokers of all ages  
#will tend to make up the majority of your expensive clients. To do risk  
#assessment, I suggest isolating the smokers and then examine them based on their  
#age and bmi. Through this you should find your expensive clientele for the year  
#to come.
```

```
#As far as how to lower your total health care costs. I think making some sort  
#of incentivisation program for the office smokers would yield the biggest total  
#cost cut. Getting even a percentage of the smoking population to quit could  
#reduce the total cost enough the make the incentives financially feasible.
```

```
#Feel free to reach out with any questions!
```

```
#thanks,
```

```
#Nick
```